


Automata-Theoretic LTL Model-Checking

Arie Gurfinkel
arie@cmu.edu

SEI/CMU

LTL - Linear Time Logic (Pn 77)

Determines Patterns on Infinite Traces



Atomic Propositions

Boolean Operations

Temporal operators

→	a	“a is true now”
	X a	“a is true in the neXt state”
	Fa	“a will be true in the F uture”
	Ga	“a will be G lobally true in the future”
	a U b	“a will hold true U ntil b becomes true”




LTL - Linear Time Logic (Pn 77)

Determines Patterns on Infinite Traces

Atomic Propositions

Boolean Operations

Temporal operators



	a	"a is true now"
→	$X a$	"a is true in the neXt state"
	Fa	"a will be true in the FUTURE"
	Ga	"a will be GLOBALLY true in the future"
	$a U b$	"a will hold true U ntil b becomes true"




LTL - Linear Time Logic (Pn 77)

Determines Patterns on Infinite Traces

Atomic Propositions

Boolean Operations

Temporal operators



a

“a is true now”

$X a$

“a is true in the neXt state”

→ $F a$

“a will be true in the Future”

$G a$

“a will be Globally true in the future”


$a U b$

“a will hold true Until b becomes true”



LTL - Linear Time Logic (Pn 77)

Determines Patterns on Infinite Traces



Atomic Propositions

Boolean Operations

Temporal operators

a "a is true now"
 $X a$ "a is true in the neXt state"
 $F a$ "a will be true in the FUTURE"


→ $G a$ "a will be Globally true in the future"

$a U b$ "a will hold true U ntil b becomes true"



LTL - Linear Time Logic (Pn 77)

Determines Patterns on Infinite Traces



Atomic Propositions

Boolean Operations

Temporal operators

a “a is true now”

X a “a is true in the ne**X**t state”

Fa “a will be true in the **F**uture”

Ga “a will be **G**lobally true in the future”

 **a U b** “a will hold true **U**ntil b becomes true”



Outline


- Automata-Theoretic Model-Checking
 - Finite Automata and Regular Languages
 - Automata over infinite words: Büchi Automata
 - Representing models and formulas with automata
 - Model checking as language emptiness

Finite Automata

A finite automaton \mathcal{A} (over finite words) is a tuple $(\Sigma, Q, \Delta, Q^0, F)$, where

- Σ is a finite alphabet
- Q is a finite set of states
- $\Delta \subseteq Q \times \Sigma \times Q$ is a transition relation
- $Q^0 \subseteq Q$ is a set of initial states
- $F \subseteq Q$ is a set of final states

Finite Automaton: An Example




$$\Sigma = \{a, b, c\}, Q = \{q_0, q_1\}, Q^0 = \{q_0\}, F = \{q_1\}$$

A Run

- A *run* of \mathcal{A} over a word $v \in \Sigma^*$ of length $|v|$ is a mapping $\rho : \{0, 1, \dots, |v|\} \rightarrow Q$ s.t.
 - First state is the initial state: $\rho(0) \in Q^0$
 - States are related by transition relation:
 $\forall 0 \leq i \leq |v| \cdot (\rho(i), v(i), \rho(i+1)) \in \Delta$
- A run is a path in \mathcal{A} from q_0 to a state $\rho(|v|)$ s.t. the edges are labeled with letters in v
- A run is *accepting* if it ends in an accepting state: $\rho(|v|) \in F$.
- \mathcal{A} *accepts* v iff exists an accepting run of \mathcal{A} on v .


An Example of a Run



- A run q_0, q_1, q_1, q_1, q_0 on $aacb$ is accepting
- A run q_0, q_0, q_0, q_0, q_0 on $bbbb$ is accepting
- A run q_0, q_0, q_1, q_1, q_1 on $baac$ is rejecting

Language

The *language* $\mathcal{L}(\mathcal{A}) \subseteq \Sigma^*$ is the set of all words in Σ^* accepted by \mathcal{A} .



The language is $\{\epsilon, b, bb, ccc, bab, \dots\}$

That is, a regular expression: $\epsilon + a(a + c)^*b(b + c)^*$

Regular Languages

- A set of strings is *regular* if it is a language of a finite automaton (i.e., recognizable by a finite automaton)
- An automaton is **deterministic** if the transition relation is deterministic for every letter in the alphabet:

$$\forall a \cdot (q, a, q') \in \Delta \wedge (q, a, q'') \in \Delta \Rightarrow q' = q''$$

otherwise, it is *non-deterministic*.

- NFA = DFA: Every non-deterministic finite automaton (NFA) can be translated into a language-equivalent deterministic automaton (DFA)

Automata on Infinite Words

- Reactive programs execute forever – need infinite sequences of states to model them!
- Solution: finite automata over infinite words.
- Simplest case: Büchi automata
 - Same structure as automata on finite words
 - ... but different notion of acceptance
 - Recognize words from Σ^ω (not Σ^* !)
 - $\Sigma = \{a, b\}$ $v = abaabaaab\dots$
 - $\Sigma = \{a, b, c\}$
 $\mathcal{L}_1 = \{v \mid \text{in } v \text{ after every } a \text{ there is a } b\}$
Some words in \mathcal{L}_1 :
 $ababab\dots$ $aaabaaab\dots$
 $abbabbabb\dots$ $accbaccb\dots$


Infinite Run and Acceptance

- Recall, F is the set of *accepting* states
- A *run* ρ of a Büchi automaton \mathcal{A} is over an infinite word $v \in \Sigma^\omega$. Domain of the run is the set of all natural numbers.
- Let $\text{inf}(\rho)$ be the set of states that appear infinitely often in ρ :


$$\text{inf}(\rho) = \{q \mid \forall i \in \mathbb{N} \cdot \exists j \geq i \cdot \rho(j) = q\}$$

- A run ρ is *accepting* (Büchi accepting) iff $\text{inf}(\rho) \cap F \neq \emptyset$.
- A set of strings is ω -regular iff it is recognizable by a Büchi automaton

Example



Example



Language of the automaton is: $((b + c)^\omega a (a + c)^* b)^\omega$

This is an ω -regular expression

Examples

Let $\Sigma = \{0, 1\}$. Define Büchi automata for the following languages:

1. $L = \{v \mid 0 \text{ occurs in } v \text{ exactly once}\}$
2. $L = \{v \mid \text{after each } 0 \text{ in } v \text{ there is a } 1\}$
3. $L = \{v \mid v \text{ contains finitely many } 1\text{'s}\}$
4. $L = (01)^n \Sigma^\omega$
5. $L = \{v \mid 0 \text{ occurs in every even position of } v\}$


Closure Properties

Büchi-recognizable languages are closed under ...

- (alphabet) projection and union
 - Same algorithms as Finite Automata
- intersection
 - Different construction from Finite Automata
- complement
 - i.e., from a Büchi automaton \mathcal{A} recognizing \mathcal{L} one can construct an automaton $\overline{\mathcal{A}}$ recognizing $\Sigma^\omega - \mathcal{L}$.
 - $\overline{\mathcal{A}}$ has order of $O(2^{Q \log Q})$ states, where Q are states in \mathcal{A} [Safra's construction]



Complementation: Example

Complement is easy for deterministic Büchi automata:





Complementation: Example

Complement is easy for deterministic Büchi automata:





Complementation: Example

Complement is easy for deterministic Büchi automata:



But, Büchi automata are not closed under determinization!!!



Intersection (Special Case)

Büchi automata are closed under intersection [Chouka74]:

- given two Büchi automata (note all states of \mathcal{B}_1 are accepting):

$$\mathcal{B}_1 = (\Sigma, Q_1, \Delta_1, Q_1^0, Q_1) \quad \mathcal{B}_2 = (\Sigma, Q_2, \Delta_2, Q_2^0, F_2)$$

- Define $\mathcal{B}_\cap = (\Sigma, Q_1 \times Q_2, \Delta', Q_1^0 \times Q_2^0, Q_1 \times F_2)$, where
 - $((s_1, s_2), a, (s'_1, s'_2)) \in \Delta'$ iff $(s_i, a, s'_i) \in \Delta_i, i = 1, 2$
- Then, $\mathcal{L}(\mathcal{B}_\cap) = \mathcal{L}(\mathcal{B}_1) \cap \mathcal{L}(\mathcal{B}_2)$

Intersection (General Case)

- Main problem: determining accepting states
 - need to go through accepting states of \mathcal{B}_1 and \mathcal{B}_2 infinite number of times

Intersection (General Case)

- Main problem: determining accepting states
 - need to go through accepting states of \mathcal{B}_1 and \mathcal{B}_2 infinite number of times
- Key idea: make 3 copies of the automaton:
 - 1st copy: start and accept here
 - 2nd copy: move from here from 1 when accepting state from \mathcal{B}_1 has been seen
 - 3rd copy: move here from 2 when accepting state from \mathcal{B}_2 has been seen, then go back to 1

Intersection (General Case)

Given two Büchi automata:

$$\mathcal{B}_1 = (\Sigma, Q_1, \Delta_1, Q_1^0, F_1) \quad \mathcal{B}_2 = (\Sigma, Q_2, \Delta_2, Q_2^0, F_2)$$

Define

$\mathcal{B}_\cap = (\Sigma, Q_1 \times Q_2 \times \{0, 1, 2\}, \Delta', Q_1^0 \times Q_2^0 \times \{0\}, Q_1 \times Q_2 \times \{2\})$,
where

- $((s_1, s_2, 0), a, (s'_1, s'_2, 0)) \in \Delta'$ iff $(s_i, a, s'_i) \in \Delta_i, i = 1, 2$, and $s'_1 \notin F_1$
- $((s_1, s_2, 1), a, (s'_1, s'_2, 1)) \in \Delta'$ iff $(s_i, a, s'_i) \in \Delta_i, i = 1, 2$, and $s'_2 \notin F_2$
- $((s_1, s_2, 0), a, (s'_1, s'_2, 1)) \in \Delta'$ iff $(s_i, a, s'_i) \in \Delta_i, i = 1, 2$, and $s'_1 \in F_1$
- $((s_1, s_2, 1), a, (s'_1, s'_2, 2)) \in \Delta'$ iff $(s_i, a, s'_i) \in \Delta_i, i = 1, 2$, and $s'_2 \in F_2$
- $((s_1, s_2, 2), a, (s'_1, s'_2, 0)) \in \Delta'$ iff $(s_i, a, s'_i) \in \Delta_i, i = 1, 2$


Then, $\mathcal{L}(\mathcal{B}_\cap) = \mathcal{L}(\mathcal{B}_1) \cap \mathcal{L}(\mathcal{B}_2)$

Complexity


- The emptiness problem for Büchi automata is decidable
 - $\mathcal{L}(\mathcal{A}) \neq \emptyset$
 - logspace-complete for NLOGSPACE, i.e., solvable in linear time [Vardi, Wolper] – see later in the lecture.
- Nonuniversality problem for Büchi automata is decidable
 - $\mathcal{L}(\mathcal{A}) \neq \Sigma^\omega$
 - logspace-complete for PSPACE [Sistla, Vardi, Wolper]

Modeling Systems Using Automata

- A system is a set of all its executions. So, every state is accepting!
- Transform Kripke structure (S, R, S_0, L) , where $L : S \rightarrow 2^{AP}$
- ...into automaton $\mathcal{A} = (\Sigma, S \cup \{\ell\}, \Delta, \{\ell\}, S \cup \{\ell\})$,
 - where $\Sigma = 2^{AP}$
 - $(\ell, \alpha, s') \in \Delta$ iff $s \in S_0$ and $\alpha = L(s)$
 - $(s, \alpha, s) \in \Delta$ iff $(s, s') \in R$ and $\alpha = L(s')$




Kripke structure



Automaton

LTL and Büchi Automata

- Specification – also in the form of an automaton!
- Büchi automata can encode all LTL properties.
- Examples:



○

$a U b$

- Other examples:
 - $\Box \Diamond p$
 - $\Box \Diamond (p \vee q)$
 - $\neg \Box \Diamond (p \vee q)$
 - $\neg(\Box(p U q))$

LTL to Büchi Automata

- Theorem [Wolper, Vardi, Sistla 83]: Given an LTL formula ϕ , one can build a Büchi automaton $\mathcal{S} = (\Sigma, Q, \Delta, Q_0, F)$, where
 - $\Sigma = 2^{\text{Prop}}$
 - the number of atomic propositions, variables, etc. in ϕ
 - $|Q| \leq 2^{O(|\phi|)}$, where $|\phi|$ is the length of the formula
- ... s.t. $\mathcal{L}(\mathcal{S})$ is exactly the set of computations satisfying the formula ϕ .
- Algorithm: see Section 9.4 of Model Checking book or try one of the online tools:
 - <http://spot.lip6.fr/ltl2tgba.html>
 - <http://www.lsv.ens-cachan.fr/~gastin/ltl2ba/index.php>
- But Büchi automata are more expressive than LTL!

Automata-theoretic Model Checking

- The system \mathcal{A} satisfies the specification \mathcal{S} when
 - $\mathcal{L}(\mathcal{A}) \subseteq \mathcal{L}(\mathcal{S})$
 - ... each behavior of the system is among the allowed behaviours
- Alternatively,
 - let $\overline{\mathcal{L}(\mathcal{S})}$ be the language $\Sigma^\omega - \mathcal{L}(\mathcal{S})$. Then,
 - $\mathcal{L}(\mathcal{A}) \subseteq \mathcal{L}(\mathcal{S}) \iff \mathcal{L}(\mathcal{A}) \cap \overline{\mathcal{L}(\mathcal{S})} = \emptyset$
 - no behavior of \mathcal{A} is prohibited by \mathcal{S}
 - If the intersection is not empty, any behavior in it corresponds to a counterexample.
 - Counterexamples are always of the form uv^ω , where u and v are finite words.

Complexity


- Checking whether a formula ϕ is satisfied by a finite-state model K can be done in time $O(\|K\| \times 2^{O(\|\phi\|)})$ or in space $O((\log\|K\| + \|\phi\|)^2)$.
- i.e., checking is polynomial in the size of the model and exponential in the size of the specification.

Emptiness of Büchi Automata

- An automaton is non-empty iff
 - there exists a path to a cycle containing an accepting state


Emptiness of Büchi Automata

- An automaton is non-empty iff
 - there exists a path to a cycle containing an accepting state
- Is this automaton empty?




Emptiness of Büchi Automata

- An automaton is non-empty iff
 - there exists a path to a cycle containing an accepting state
- Is this automaton empty?
 - No – it accepts $a(bef)^\omega$




Emptiness of Büchi Automata

- An automaton is non-empty iff
 - there exists a path to a cycle containing an accepting state
- Is this automaton empty?
 - No – it accepts $a(bef)^\omega$




Emptiness of Büchi Automata

- An automaton is non-empty iff
 - there exists a path to a cycle containing an accepting state
- Is this automaton empty?
 - No – it accepts $a(bef)^\omega$




Emptiness of Büchi Automata

- An automaton is non-empty iff
 - there exists a path to a cycle containing an accepting state
- Is this automaton empty?
 - No – it accepts $a(bef)^\omega$



Emptiness of Büchi Automata


- An automaton is non-empty iff
 - there exists a path to a cycle containing an accepting state
- Is this automaton empty?
 - No – it accepts $a(bef)^\omega$




LTL Model-Checking

- LTL Model-Checking = Emptiness of Büchi automata
 - a tiny bit of automata theory +
 - trivial graph-theoretic problem
 - typical solution – use depth-first search (DFS)
- Problem: **state-explosion**
 - the graph is *HUGE*
- End result:
 - LTL model-checking a very elaborate DFS


Depth-First Search – Refresher




Depth-First Search – Refresher




Depth-First Search – Refresher




Depth-First Search – Refresher




Depth-First Search – Refresher




Depth-First Search – Refresher




Depth-First Search – Refresher




Depth-First Search – Refresher



Depth-First Search – Refresher



Depth-first tree



DFS – The Algorithm

```
1: time := 0
2: proc DFS(v)
3:   add v to Visited
4:   d[v] := time
5:   time := time + 1
6:   for all w ∈ succ(v) do
7:     if w ∉ Visited then
8:       DFS(w)
9:     end if
10:  end for
11:  f[v] := time
12:  time := time + 1
13: end proc
```

DFS – Data Structures

- implicit STACK
 - stores the current path through the graph
- *Visited* table
 - stores visited nodes
 - used to avoid cycles
- for each node
 - *discovery time* – array d
 - *finishing time* – array f

What we want

- Running time
 - at most linear — anything else is not feasible
- Memory requirements
 - sequentially accessed – (for the STACK)
 - disk storage is good enough
 - assume unlimited supply – so can ignore
 - randomly accessed – (for hash tables)
 - must use RAM
 - limited resource – minimize
 - why cannot use virtual memory?

Additionally...

- Counterexamples
 - an automaton is non-empty iff exists an accepting run
 - this is the counterexample – we want it
- Approximate solutions
 - partial result is better than nothing!

DFS – Complexity

- Running time
 - each node is visited once
 - linear in the size of the graph
- Memory
 - the STACK
 - accessed sequentially
 - can store on disk – ignore
 - *Visited* table
 - randomly accessed – important
 - $|Visited| = S \times n$
 - n – number of nodes in the graph
 - S – number of bits needed to represent each node

Take 1 – Tarjan's SCC algorithm



- Idea: find all maximal SCCs: SCC_1 , SCC_2 , etc.
 - an automaton is non-empty iff exists SCC_i containing an accepting state

Take 1 – Tarjan's SCC algorithm

- Idea: find all maximal SCCs: SCC_1 , SCC_2 , etc.
 - an automaton is non-empty iff exists SCC_i containing an accepting state
- Fact: each SCC is a sub-tree of DFS-tree
 - need to find roots of these sub-trees

Take 1 – Tarjan's SCC algorithm

- Idea: find all maximal SCCs: SCC_1 , SCC_2 , etc.
 - an automaton is non-empty iff exists SCC_i containing an accepting state
- Fact: each SCC is a sub-tree of DFS-tree
 - need to find roots of these sub-trees



Finding a Root of an SCC


- For each node v , compute $lowlink[v]$
- $lowlink[v]$ is the minimum of
 - discovery time of v
 - discovery time of w , where
 - w belongs to the same SCC as v
 - the length of a path from v to w is at least 1
- Fact: v is a root of an SCC iff
 - $d[v] = lowlink[v]$

Finally: the algorithm

```
1: proc SCC_SEARCH(v)
2:   add v to Visited
3:   d[v] := time
4:   time := time + 1
5:   lowlink[v] := d[v]
6:   push v on STACK
7:   for all w ∈ succ(v) do
8:     if w ∉ Visited then
9:       SCC_SEARCH(w)
10:    lowlink[v] := min(lowlink[v], lowlink[w])
11:    else if d[w] < d[v] and w is on STACK then
12:      lowlink[v] := min(d[w], lowlink[v])
13:    end if
14:  end for
15:  if lowlink[v] = d[v] then
16:    repeat
17:      pop x from top of STACK
18:      if x ∈ F then
19:        terminate with “Yes”
20:      end if
21:    until x = v
22:  end if
23: end proc
```

Finally: the algorithm

```
1: proc SCC_SEARCH(v)
2:   add v to Visited
3:   d[v] := time
4:   time := time + 1
5:   lowlink[v] := d[v]
6:   push v on STACK
7:   for all w ∈ succ(v) do
8:     if w ∉ Visited then
9:       SCC_SEARCH(w)
10:    lowlink[v] := min(lowlink[v], lowlink[w])
11:  else if d[w] < d[v] and w is on STACK then
12:    lowlink[v] := min(d[w], lowlink[v])
13:  end if
14: end for
15: if lowlink[v] = d[v] then
16:   repeat
17:     pop x from top of STACK
18:     if x ∈ F then
19:       terminate with “Yes”
20:     end if
21:   until x = v
22: end if
23: end proc
```



Tarjan's SCC algorithm – Analysis

- Running time
 - linear in the size of the graph
- Memory
 - STACK – sequential, ignore
 - *Visited* – $O(S \times n)$
 - *lowlink* – $\log n \times n$
 - n is not known a priori
 - assume n is at least $\geq 2^{32}$
- Counterexamples
 - can be extracted from the STACK
 - even more – get multiple counterexamples
- If we sacrifice some of generality, can we do better?

Take 2 – Two Sweeps


- Don't look for maximal SCCs
- Find a reachable accepting state that is on a cycle
- Idea: use two sweeps
 - sweep one: find all accepting states
 - sweep two: look for cycles *from* accepting states

Take 2 – Two Sweeps

- Don't look for maximal SCCs
- Find a reachable accepting state that is on a cycle
- Idea: use two sweeps
 - sweep one: find all accepting states
 - sweep two: look for cycles *from* accepting states
- Problem?
 - no longer a linear algorithm (revisit the states multiple times)

Take 2 – Two Sweeps

- Don't look for maximal SCCs
- Find a reachable accepting state that is on a cycle
- Idea: use two sweeps
 - sweep one: find all accepting states
 - sweep two: look for cycles *from* accepting states
- Problem?
 - no longer a linear algorithm (revisit the states multiple times)




Fixing non-linearity

- Graph Theoretic Result: let v and u be two nodes, such that
 - $f[v] < f[u]$
 - v is not on a cycle
 - then, no cycle containing u contains nodes reachable from v


Fixing non-linearity

- Graph Theoretic Result: let v and u be two nodes, such that
 - $f[v] < f[u]$
 - v is not on a cycle
 - then, no cycle containing u contains nodes reachable from v



Fixing non-linearity

- Graph Theoretic Result: let v and u be two nodes, such that
 - $f[v] < f[u]$
 - v is not on a cycle
 - then, no cycle containing u contains nodes reachable from v



Take 3 – Double DFS

```
1: proc DFS1(v)
2:   add v to Visited
3:   for all w ∈ succ(v) do
4:     if w ∉ Visited then
5:       DFS1(w)
6:     end if
7:   end for
8:   if v ∈ F then
9:     add v to Q
10:  end if
11: end proc
```

```
1: proc DFS2(v, f)
2:   add v to Visited
3:   for all w ∈ succ(v) do
4:     if v = f then
5:       terminate with “Yes”
6:     else if w ∉ Visited then
7:       DFS2(w, f)
8:     end if
9:   end for
10: end proc
```

```
1: proc SWEEP2(Q)
2:   while Q ≠ [] do
3:     f := dequeue(Q)
4:     DFS2(f, f)
5:   end while
6:   terminate with “No”
7: end proc
```

```
1: proc DDFS(v)
2:   Q = ∅
3:   Visited = ∅
4:   DFS1(v)
5:   Visited = ∅
6:   SWEEP2(Q)
7: end proc
```

Double DFS – Analysis

- Running time
 - linear! (single *Visited* table for different final states, so no state is processed twice)
- Memory requirements
 - $O(n \times S)$
- Problem
 - where is the counterexample?!

Take 4 – Nested DFS

- Idea
 - when an accepting state is finished
 - stop first sweep
 - start second sweep
 - if cycle is found, we are done
 - otherwise, restart the first sweep
- As good as double DFS, but
 - does not need to *always* explore the full graph
 - counterexample is readily available
 - a path to an accepting state is on the stack of the first sweep
 - a cycle is on the stack of the second

A Few More Tweaks

- No need for two *Visited* hashtables
 - empty hashtable wastes space
 - merge into one by adding one more bit to each node
 - $(v, 0) \in Visited$ iff v was seen by the first sweep
 - $(v, 1) \in Visited$ iff v was seen by the second sweep
- Early termination condition
 - nested DFS can be terminated as soon as it finds a node that is on the stack of the first DFS

Nested DFS

```
1: proc DFS1(v)
2:   add (v, 0) to Visited
3:   for all w ∈ succ(v) do
4:     if (w, 0) ∉ Visited then
5:       DFS1(w)
6:     end if
7:   end for
8:   if v ∈ F then
9:     DFS2(v, v)
10:  end if
11: end proc
```

```
1: proc DFS2(v, f)
2:   add (v, 1) to Visited
3:   for all w ∈ succ(v) do
4:     if v = f then
5:       terminate with “Yes”
6:     else if (w, 1) ∉ Visited then
7:       DFS2(w, f)
8:     end if
9:   end for
10: end proc
```


On-the-fly Model-Checking

- Typical problem consists of
 - description of several process P_1, P_2, \dots
 - property φ in LTL
- Before applying DFS algorithm
 - construct graph for $P = \prod_{i=1}^n P_i$
 - construct Büchi automaton $A_{\neg\varphi}$ for $\neg\varphi$
 - construct Büchi automaton for $P \cap A_{\neg\varphi}$

On-the-fly Model-Checking

- Typical problem consists of
 - description of several process P_1, P_2, \dots
 - property φ in LTL
- Before applying DFS algorithm
 - construct graph for $P = \prod_{i=1}^n P_i$
 - construct Büchi automaton $A_{\neg\varphi}$ for $\neg\varphi$
 - construct Büchi automaton for $P \cap A_{\neg\varphi}$
- But,
 - all constructions can be done in DFS order
 - combine everything with the search
 - result: on-the-fly algorithm, only the necessary part of the graph is built

Symbolic LTL Model-Checking

- LTL Model-Checking = Finding a reachable cycle
- Represent the graph symbolically
 - and use symbolic techniques to search
- There exists an infinite path from s , iff $s \models EG \text{ true}$
 - the graph is finite
 - infinite \Rightarrow cyclic!
 - exists a cycle containing an accepting state a iff a occurs infinitely often
 - use fairness to capture accepting states
- LTL Model-Checking = $EG \text{ true}$ under fairness!